

Readme for the Self-Driving Cars Nanodegree Program 3rd Project

Traffic Sign Recognition

Goals / Steps of this project:

- Load the data set.
- Explore, summarize and visualize the data set.
- Design, train and test a model architecture.
- Use the model to make predictions on new images.
- Analyze the softmax probabilities of the new images.
- Summarize the results with a written report.

Rubric Points

Dataset Exploration

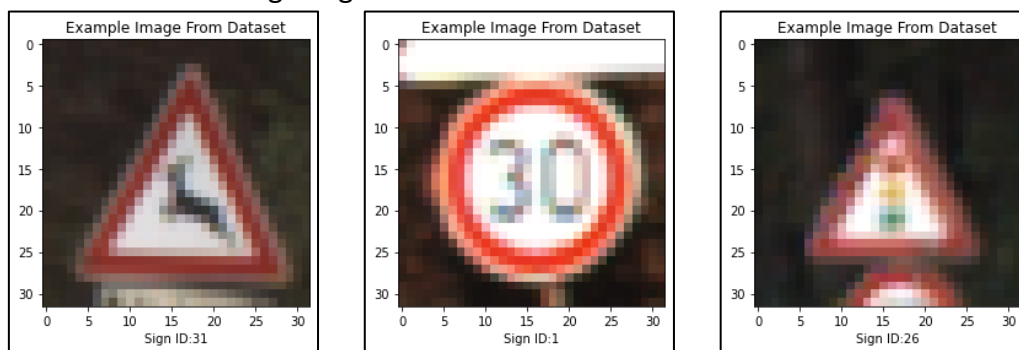
Dataset Summary

- German Traffic Sign Dataset was used to train the model used for the Traffic Sign Recognition.
 - Size of **Training** set is: 34,799
 - Size of **Validation** set is: 4,410
 - Size of **Testing** set is: 12630
 - Image data shape is: (32, 32, 3)
 - Number of Classes: 43*


*The different Classes identification can be seen in the "signnames.csv" included in the project delivery.


Exploratory Visualization

- The Following images were taken from the dataset.



- In the following charts is shown how the training, validation and testing data was divided.






Design and Test a Model Architecture



Pre-processing

- The pre-processing of the data was done in 4 steps:
 - Shuffle the training data: This step is done to increase randomness and variety in the training dataset.
 - Convert Images to Grayscale: This step is done to all the input data. The use of grayscale images instead of colour ones improves the Convnets accuracy.
 - Apply Local Histogram Equalization: This step increases the contrast of the images. This is helpful since the images are taken from real world and many of them have low contrast.
 - Normalize the Data: Once the Images are converted to grayscale and the depth is just 1, I normalize the image by dividing every pixel [0,255] by 255; the result is a value between [0,1].

Note: In the code a function to increase the dataset can be found "add_rotate". This Functions takes random images from the dataset and randomly rotates it and adds it again to the dataset. Finally, this function was not used since during the test I saw it did not contribute too much to the CNN Accuracy.

- The following images are examples of Traffic signs before/after being pre-processed.





Model Architecture

- The table describes the layers of the Model.

Layer	Description
Input	Grayscale Image 32x32x1
Convolution	Stride: 1x1, Padding: Valid, Output: 28x28x6
ReLU	
Max Pooling	Stride: 2x2, Padding: Valid, Output: 14x14x6
Convolution	Stride: 1x1, Padding: Valid, Output: 10x10x16
ReLU	
Max Pooling	Stride: 2x2, Padding: Valid, Output: 5x5x16
Flatten	Output: 400
Fully-Connected	Output: 120
ReLU	
Dropout	Training keep_prob: 0.5
Fully-Connected	Output: 84
ReLU	
Dropout	Training Keep_prob: 0.5
Fully-Connected	Output: 43

Model Training

- Hyperparameters
 - Epochs: 30
 - Batch_size: 128
 - Learning rate: 0.001
- To train the model, the the LeNet CNN was used. Modifications in the output layer; and dropouts were added in the fully connected layers to avoid overfitting since it is a serious problem in deep neural networks.
 - The output was modified to give an output of 43 classes instead of ten as used during the exercises of the course.
- The final model results were the following:
 - Training set accuracy: 0.992
 - Validation set accuracy: 0.951
 - Testing set accuracy: 0.933

- The following piece of code was used to train the network:

```
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    num_examples = len(X_train)

    print("Training...")
    print()
    for i in range(EPOCHS):
        X_train, y_train = shuffle(X_train, y_train)
        for offset in range(0, num_examples, BATCH_SIZE):
            end = offset + BATCH_SIZE
            batch_x, batch_y = X_train[offset:end], y_train[offset:end]
            sess.run(training_operation, feed_dict={x: batch_x, y: batch_y, keep_prob: 0.5})

        training_accuracy = evaluate(X_train, y_train)
        validation_accuracy = evaluate(X_valid, y_valid)
        print("EPOCH {} ...".format(i+1))
        print("Training Accuracy = {:.3f}".format(training_accuracy))
        print("Validation Accuracy = {:.3f}".format(validation_accuracy))
        print()

    saver.save(sess, './lenet')
    print("Model saved")
```

Note: During training the values for Dropouts was keep_prob = 0.5 while for the evaluation was 1.0.

```
###Model Evaluation
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_y, 1))
accuracy_operation = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):
    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, BATCH_SIZE):
        batch_x, batch_y = X_data[offset:offset+BATCH_SIZE], y_data[offset:offset+BATCH_SIZE]
        accuracy = sess.run(accuracy_operation, feed_dict={x: batch_x, y: batch_y, keep_prob: 1.0})
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples
```

Solution Approach

- The first step I took was to reuse the LeNet Model designed during the course.
 - Modifications to the input and output layers were done to fit the input images (RGB) and the output (43 Classes).
 - The Validation Accuracy was around 87%, not so far from the requested Accuracy that was 93%.
- During the process to get the requested accuracy I tried many things.
 - Normalize the data with the following formula $(\text{pixel}-128)/128$ getting pixels with values from $[-1,1]$.
 - Using Sigmoid Activations instead of RELU; I tried this because of the normalized data $[-1,1]$. RELU might not be the best activation if the inputs can be negative.
 - Augment data to the dataset with the same images but randomly rotated.
 - Grayscale input.
 - Add Dropouts to all the layers (including Convolutions).
- Some of the accuracies that I got are:
 - Just adding Dropouts to all the layers: 0.637
 - Normalize dataset, Add Dropouts and augment 20% the dataset: 0.654

- Grayscale the input: 0.895
- Grayscale, Normalize data [0,1]: 0.911
- **Grayscale, Normalize data [0,1], Dropout (Just Fully-connected layers): 0.930**
 - This was the first try that fulfilled the requested Validation Accuracy.
- **Grayscale and Local Equalization, Normalize data [0,1], Dropout (Just Fully-connected layers): 0.951**
 - This one was the selected architecture (previously explained).


Note: I tried around 20 different models before obtaining a good result.

- During this tries I made a deep research to find what could improve my model, two tips that I consider very valuables from this research are:
 - When a CNN is overfitting, Dropouts are a good option to apply but just in the fully-connected layers.
 - CNN responds better to Grayscale images than Colour ones.

Test a Model on New Images

Acquiring New Images

- The selected images from the web are shown below:



- Before testing the model with the previous images, I thought it would be difficult to correctly classify the "speed limitation" and the "No Passing" images.
 - The 100km/h limitation sign because it is very similar to all the signs that limit the speed.
 - The No Passing sign because the background of the picture is not very clean.

Performance on New Images

- The model was able to correctly guess the 5 traffic signs, which gives an accuracy of 100%.
 - The Test set accuracy was: 93.3%

Model Certainty – Softmax Probabilities

- The top five softmax probabilities of the predictions on the New Images are shown below.
 - For 4 out of 5 images, the model is relatively sure of the sign, as guessed the “No Passing” Sign makes the model wonder a little bit since there are other traffic signs that are similar, like “End of no Passing” and “No Passing for vehicles over 3.5 metric tons”.

